




Subroutines & Operating Systems

Part 8

1




The System Stack

Pile of... Data

2

The System Stack

- The processor maintains a stack in memory
- It allows *subroutines*
 - analogous to the "functions" you use in Java and other third-generation languages
 - but, much more simple



Spring 2022

Subroutines: Stack - Cook - CS50.10

3

3

Examples of Stacks

- Page-visited "back button" history in a web browser
- Undo sequence in a text editor
- Deck of cards in Windows Solitaire



Spring 2022

Subroutines: Stack - Cook - CS50.10

4

4

Implementing in Memory

- On a processor, the stack stores integers
 - size of the integer the bit-size of the system
 - 64-bit system → 64-bit integer
- Stacks is stored in memory
 - A fixed location pointer (S0) defines the bottom of the stack
 - A *stack pointer* (SP) gives the location of the top of the stack

Spring 2022

Subroutines: Stack - Cook - CS50.10

5

5

Approaches

- Growing upwards
 - Bottom Pointer (S0) is the *lowest* address in the stack buffer
 - stack grows towards *higher* addresses
- Grow downwards
 - Bottom Pointer (S0) is the *highest* address in the stack buffer
 - stack grows towards *lower* addresses

Spring 2022

Subroutines: Stack - Cook - CS50.10

6

6

Size of the Stack

- As an abstract data structure...
 - stacks are assumed to be **infinitely** deep
 - so, an arbitrary amount of data can be stored
- However...
 - stacks are implemented using memory buffers
 - which are **finite** in size
- If the data **exceeds** the allocated space, a **stack overflow** error occurs

Spring 2022

Subroutines: Stack - Cook - CS30.35

7

7

Subroutine Call Basics

Organizing Your Program

8

Subroutine Call

- The stack is essential for subroutines to work
- How?
 - used to save the return addresses for call instructions
 - backup and restore registers
 - pass data between subroutines



Spring 2022

Subroutines: Stack - Cook - CS30.35

9

9

When you call a subroutine...

1. Processor pushes the instruction pointer (IP) – an address – on the stack
2. IP is set to the address of the subroutine
3. Subroutine executes and ends with a "return" instruction
4. Processor pops & restores the original IP
5. Execution continues after the initial call

Spring 2022


Subroutines: Stack - Cook - CS30.35

10

10

Nesting is Possible

- Subroutines can call other subroutines
- f()** calls **g()** which then calls **h()**, etc...
- The stack stores the return addresses of the callers
- Just like the "history button" in your web browser, you can store many return addresses



Spring 2022


Subroutines: Stack - Cook - CS30.35

11

11

Nesting is Possible

- Each time a subroutine completes, the processor pops the top of the stack
- ...then returns to the **caller**
- This allows normal function calls and recursion (a powerful tool)




Spring 2022

Subroutines: Stack - Cook - CS30.35

12

12




x64 Subroutines

Organizing Your Programs ... with Intel

13

Instruction: Call

- The *Call Instruction* transfers control to a subroutine
- Other processors call it different names such as JSR (Jump Subroutine)
- The stack is used to save the current IP



Spring 2022 Sacramento State - CS&E - CS&E 35 14

14

Instruction: Call

Usually, a label
(which is an address)

CALL *address*

Spring 2022 Sacramento State - CS&E - CS&E 35 15

15

Instruction: Return

- The Return Instruction is used mark the end of subroutine
- When the instruction is executed...
 - the old instruction pointer is read from the system stack
 - the current instruction pointer is updated – restoring execution after the initial call

Spring 2022 Sacramento State - CS&E - CS&E 35 16

16

Instruction: Return

- Do not forget this!
- If you do...
 - execution will simply continue, in memory, until a return instruction is encountered
 - often is can run past the end of your program
 - ...and run data!

Spring 2022 Sacramento State - CS&E - CS&E 35 17

17

Instruction: Return

No arguments!

RET

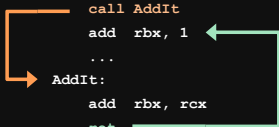
Spring 2022 Sacramento State - CS&E - CS&E 35 18

18

Subroutine Example

```

_start:
    mov rcx, 4
    mov rbx, 12
    call AddIt
    add rbx, 1
    ...
AddIt:
    add rbx, rcx
    ret
    
```



19




Operating Systems

The master software

20

What is an operating system?


- The operating system is simply a series of programs
- These programs, however, run with special privileges which are needed by the OS
- Processors support two modes for executing programs



21

Execution Modes


- **Privileged (supervisor) mode**
 - can run special instructions
 - can talk to all the hardware
 - etc...
- **User mode**
 - can only execute certain instructions
 - can't talk to all the hardware



22

Vector Tables


- Programs (and hardware) often need to talk to the operating system
- Examples:
 - software needs talk to the OS
 - USB port notifies the OS that a device was plugged in



23

Vector Tables

- But how does this happen?
- The processor can be **interrupted** – **alerted** – that something must be handled
- It then runs a special program that handles the event



24

Vector Table



- During an interrupt, the device sends the processor an *interrupt number*
- The processor looks up the number in the *vector table*
- Table contains the address of *Interrupt Service Routine (ISR)* to execute


Spring 2022

Background: State - Cook - CSU 35

25

25

How It Works



Spring 2022

Background: State - Cook - CSU 35

26

26

The Processor Actions



1. Backup the register file
2. Execute *Interrupt Service Routine (ISR)*
3. Once completed: restore the original executing program & register file

Spring 2022


Background: State - Cook - CSU 35

27

27

The Kernal

- All these Interrupt Service Routines belong to the *kernal* – the core of the operating system
- Vast majority of the operating system is hidden from the end user



Spring 2022

Background: State - Cook - CSU 35

28

28

Interact with Applications



How do WE talk to the OS

Interact with Applications

- Software also needs to talk to the operating system
- For example:
 - draw a button
 - print a document
 - close this program
 - etc...



Spring 2022

Background: State - Cook - CSU 35

30

30

Interact with Applications

- Software can interrupt itself with a specific number
- This interrupt is *designated specifically for software*
- The operating system then handles the software's request



Spring 2022

Background: State - Cook - CSU 38

31

31

Application Program Interface

- Programs "talk" to the OS using Application Program Interface (API)
- Application → Operating System → IO
- Benefits:
 - makes applications faster and smaller
 - also makes the system more secure since apps do not directly talk to IO

Spring 2022

Background: State - Cook - CSU 35

32

32

Instruction: syscall (64-bit)

SYSCALL

Calls interrupt number reserved for programs needing attention

Spring 2022

Background: State - Cook - CSU 35

33

33

Subroutine vs. Interrupt

| Subroutine | Interrupt |
|---------------------------|---------------------------|
| Executes code | Executes code |
| Returns when complete | Returns when complete |
| Called by the application | Executed by the processor |
| Part of the application | Handles events for the OS |

Spring 2022

Background: State - Cook - CSU 35

34

34

Linux System Calls

How software and hardware "talk"



Spring 2022


Background: State - Cook - CSU 35

35

35

Interrupts on the Linux

- Linux, like other operating systems communicate with applications using interrupts
- Applications do not know where (in memory) to contact the kernel – so they ask the processor to do it



Spring 2022


Background: State - Cook - CSU 35

36

36

How It Works

1. Fill the registers
2. Interrupt using `syscall`
(or `INT 0x80` if on 32-bit)
3. Any results will be stored in the registers




Spring 2022

Backwards State - Cook - CSU 35

37

37

How to Call Linux – 64 bit



- The **rax** register must contain the *system call number*
- This number indicates what you asking the OS to do
- There are only **329** total calls in the entire 64-bit UNIX operating system!


Spring 2022

Backwards State - Cook - CSU 35

38

38

How to Call Linux – 64 bit



- Different registers are used to hold data
- The order is also quite odd:
rdi, rsi, rdx, r10, r8

Spring 2022


Backwards State - Cook - CSU 35

39

39

Kernels are Simple!

- Linux only has **1** write and **1** read system call
- The location, number of bytes, and device only change
"write x many bytes from address y to device z"
- So, writing to the screen, a file, a port, etc...use the **same** call!



Spring 2022

Backwards State - Cook - CSU 35

40

40

Some Linux 64 Calls

| System Call | rax | rdi | rsi | rdx |
|-------------|-----|-----------------|---------|-----------|
| read | 0 | file descriptor | address | max bytes |
| write | 1 | file descriptor | address | count |
| open | 2 | address | flags | mode |
| close | 3 | file descriptor | | |
| get pid | 39 | | | |
| exit | 60 | error code | | |

Spring 2022

Backwards State - Cook - CSU 35

41

41

Linux 64: Sys Write

```

mov rax, 1
mov rdi, 1
lea rsi, address
mov rdx, length
syscall
    
```

Linux command for WRITE

1 = Screen

Call Linux

Spring 2022

Backwards State - Cook - CSU 35

42

42

Linux 64: Sys Read

```
mov rax, 0
mov rdi, 0
lea rsi, address
mov rdx, maxBytes
syscall
```

Linux command for READ

0 = Keyboard

Maximum number of bytes to read

Call Linux

43

Saving Registers & Lost Data

Avoiding horrible side-effects

44

Saving Registers & Lost Data

- Each subroutine will use the registers as it needs
- So, when a sub is called, *it may modify the caller's registers*
- Some processors have few registers – so its *very* likely
- This can lead to hard-to-fix bugs if caution is not used – e.g. loop counter gets changed



45

Two Solutions

- Caller saves values
 - caller saves *all* their registers to memory before making the subroutine call
 - after, it restores the values before continuing
 - not recursion friendly – it pushes all of them!
- Subroutine saves the values
 - push registers (it will change) onto the stack
 - before it returns, it pops (and restores) the old values off the stack

46

Saving Registers... How nice! :-)

```
DoSomething:
push rax
push rbx
push rcx

...

pop rcx
pop rbx
pop rax
ret
```

Save registers

Your code

Restore them.
Note the reverse order

47